

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [4]: dataset=pd.read_csv('loan.csv')
```


```
In [5]: dataset
```

Out[5]:


	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan
	0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN
	1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0
	2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0
	3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0
	4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0

	609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0
	610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0
	611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0
	612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0
	613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0

614 rows × 13 columns



```
In [9]: sns.heatmap(dataset.isnull(),yticklabels=False,cbar=False,cmap='viridis')
plt.show()
```



```
In [11]: dataset.describe()
```

Out[11]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.000000	0.000000
25%	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	3812.500000	1188.500000	128.000000	360.000000	1.000000
75%	5795.000000	2297.250000	168.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

In [13]:

```
dataset.isnull().sum()
```

Out[13]:

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
dtype:	int64

In [15]:

```
pd.crosstab(dataset['Credit_History'], dataset['Loan_Status'], margins=True)
```

Out[15]:

Loan_Status	N	Y	All
Credit_History			
0.0	82	7	89
1.0	97	378	475
All	179	385	564

In [17]:

```
dataset['Gender'].fillna(dataset['Gender'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\2782724874.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

dataset['Gender'].fillna(dataset['Gender'].mode()[0],inplace=True)

In [19]:

```
dataset['Married'].fillna(dataset['Married'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\568568310.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

dataset['Married'].fillna(dataset['Married'].mode()[0],inplace=True)

In [21]:

```
dataset['Self_Employed'].fillna(dataset['Self_Employed'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\39551059.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

dataset['Self_Employed'].fillna(dataset['Self_Employed'].mode()[0],inplace=True)

```
In [23]: dataset['Dependents'].fillna(dataset['Dependents'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\2291896846.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Dependents'].fillna(dataset['Dependents'].mode()[0],inplace=True)
```

```
In [25]: dataset['Loan_Amount_Term'].fillna(dataset['Loan_Amount_Term'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\3879141361.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Loan_Amount_Term'].fillna(dataset['Loan_Amount_Term'].mode()[0],inplace=True)
```

```
In [27]: dataset['Credit_History'].fillna(dataset['Credit_History'].mode()[0],inplace=True)
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel_12460\979352118.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
dataset['Credit_History'].fillna(dataset['Credit_History'].mode()[0],inplace=True)
```

```
In [29]: dataset.LoanAmount=dataset.LoanAmount.fillna(dataset.LoanAmount.mean())
```

```
In [31]: dataset.isnull().sum()
```

```
Out[31]: Loan_ID          0
Gender          0
Married         0
Dependents      0
Education       0
Self_Employed   0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      0
Loan_Amount_Term 0
Credit_History  0
Property_Area   0
Loan_Status     0
dtype: int64
```

```
In [33]: dataset
```

Out[33]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan
0	LP001002	Male	No	0	Graduate	No	5849	0.0	146.412162	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.000000	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.000000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.000000	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.000000	
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.000000	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.000000	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.000000	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.000000	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.000000	

614 rows × 13 columns



In [35]:

```
gender=pd.get_dummies(dataset['Gender'],drop_first=True)
married=pd.get_dummies(dataset['Married'],drop_first=True)
education=pd.get_dummies(dataset['Education'],drop_first=True)
self_employed=pd.get_dummies(dataset['Self_Employed'],drop_first=True)
dependents=pd.get_dummies(dataset['Dependents'],drop_first=True)
property_area=pd.get_dummies(dataset['Property_Area'],drop_first=True)
dataset.drop(['Gender','Married','Education','Self_Employed','Dependents','Property_Area'],axis=1,inplace=True)
```

In [37]:

```
dataset=pd.concat([dataset,gender,married,education,self_employed,dependents,property_area],axis=1)
```

In [39]:

```
dataset
```

Out[39]:

	Loan_ID	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Male	Yes
0	LP001002	5849	0.0	146.412162	360.0	1.0	Y	True	False
1	LP001003	4583	1508.0	128.000000	360.0	1.0	N	True	True
2	LP001005	3000	0.0	66.000000	360.0	1.0	Y	True	True
3	LP001006	2583	2358.0	120.000000	360.0	1.0	Y	True	True
4	LP001008	6000	0.0	141.000000	360.0	1.0	Y	True	False
...
609	LP002978	2900	0.0	71.000000	360.0	1.0	Y	False	False
610	LP002979	4106	0.0	40.000000	180.0	1.0	Y	True	True
611	LP002983	8072	240.0	253.000000	360.0	1.0	Y	True	True
612	LP002984	7583	0.0	187.000000	360.0	1.0	Y	True	True
613	LP002990	4583	0.0	133.000000	360.0	0.0	N	False	False

614 rows × 16 columns



In [41]:

```
dataset['Male']=dataset['Male'].astype(int)
dataset['Yes']=dataset['Yes'].astype(int)
dataset['Not Graduate']=dataset['Not Graduate'].astype(int)
```

In [43]:

```
dataset
```

	Loan_ID	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Male	Yes	G
0	LP001002	5849	0.0	146.412162	360.0	1.0	Y	1	0	
1	LP001003	4583	1508.0	128.000000	360.0	1.0	N	1	1	
2	LP001005	3000	0.0	66.000000	360.0	1.0	Y	1	1	
3	LP001006	2583	2358.0	120.000000	360.0	1.0	Y	1	1	
4	LP001008	6000	0.0	141.000000	360.0	1.0	Y	1	0	
...	
609	LP002978	2900	0.0	71.000000	360.0	1.0	Y	0	0	
610	LP002979	4106	0.0	40.000000	180.0	1.0	Y	1	1	
611	LP002983	8072	240.0	253.000000	360.0	1.0	Y	1	1	
612	LP002984	7583	0.0	187.000000	360.0	1.0	Y	1	1	
613	LP002990	4583	0.0	133.000000	360.0	0.0	N	0	0	

```
dataset['Yes']=dataset['Yes'].astype(int)
dataset['1']=dataset['1'].astype(int)
dataset['2']=dataset['2'].astype(int)
dataset['3+']=dataset['3+'].astype(int)
dataset['Semiurban']=dataset['Semiurban'].astype(int)
dataset['Urban']=dataset['Urban'].astype(int)
```

dataset

	Loan_ID	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Male	Yes	G
0	LP001002	5849	0.0	146.412162	360.0	1.0	Y	1	0	
1	LP001003	4583	1508.0	128.000000	360.0	1.0	N	1	1	
2	LP001005	3000	0.0	66.000000	360.0	1.0	Y	1	1	
3	LP001006	2583	2358.0	120.000000	360.0	1.0	Y	1	1	
4	LP001008	6000	0.0	141.000000	360.0	1.0	Y	1	0	
...	
609	LP002978	2900	0.0	71.000000	360.0	1.0	Y	0	0	
610	LP002979	4106	0.0	40.000000	180.0	1.0	Y	1	1	
611	LP002983	8072	240.0	253.000000	360.0	1.0	Y	1	1	
612	LP002984	7583	0.0	187.000000	360.0	1.0	Y	1	1	
613	LP002990	4583	0.0	133.000000	360.0	0.0	N	0	0	

```
loan_status=pd.get_dummies(dataset['Loan_Status'],drop_first=True)
dataset.drop(['Loan_Status'],axis=1,inplace=True)
```

```
dataset=pd.concat([dataset,loan_status],axis=1)
```

dataset

	Loan_ID	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Male	Yes	Not Graduate	Yes
0	LP001002	5849	0.0	146.412162	360.0	1.0	1	0	0	0
1	LP001003	4583	1508.0	128.000000	360.0	1.0	1	1	0	0
2	LP001005	3000	0.0	66.000000	360.0	1.0	1	1	0	1
3	LP001006	2583	2358.0	120.000000	360.0	1.0	1	1	1	0
4	LP001008	6000	0.0	141.000000	360.0	1.0	1	0	0	0
...
609	LP002978	2900	0.0	71.000000	360.0	1.0	0	0	0	0
610	LP002979	4106	0.0	40.000000	180.0	1.0	1	1	0	0
611	LP002983	8072	240.0	253.000000	360.0	1.0	1	1	0	0
612	LP002984	7583	0.0	187.000000	360.0	1.0	1	1	0	0
613	LP002990	4583	0.0	133.000000	360.0	0.0	0	0	0	1

```
dataset['Y']=dataset['Y'].astype(int)
```

dataset

	Loan_ID	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Male	Yes	Not Graduate	Yes
0	LP001002	5849	0.0	146.412162	360.0	1.0	1	0	0	0
1	LP001003	4583	1508.0	128.000000	360.0	1.0	1	1	0	0
2	LP001005	3000	0.0	66.000000	360.0	1.0	1	1	0	1
3	LP001006	2583	2358.0	120.000000	360.0	1.0	1	1	1	0
4	LP001008	6000	0.0	141.000000	360.0	1.0	1	0	0	0
...
609	LP002978	2900	0.0	71.000000	360.0	1.0	0	0	0	0
610	LP002979	4106	0.0	40.000000	180.0	1.0	1	1	0	0
611	LP002983	8072	240.0	253.000000	360.0	1.0	1	1	0	0
612	LP002984	7583	0.0	187.000000	360.0	1.0	1	1	0	0
613	LP002990	4583	0.0	133.000000	360.0	0.0	0	0	0	1

```
dataset.drop('Loan_ID',axis=1,inplace=True)
```

dataset

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Male	Yes	Not Graduate	Yes	1	2	3+
0	5849	0.0	146.412162	360.0	1.0	1	0	0	0	0	0	0
1	4583	1508.0	128.000000	360.0	1.0	1	1	0	0	1	0	0
2	3000	0.0	66.000000	360.0	1.0	1	1	0	1	0	0	0
3	2583	2358.0	120.000000	360.0	1.0	1	1	1	0	0	0	0
4	6000	0.0	141.000000	360.0	1.0	1	0	0	0	0	0	0
...
609	2900	0.0	71.000000	360.0	1.0	0	0	0	0	0	0	0
610	4106	0.0	40.000000	180.0	1.0	1	1	0	0	0	0	1
611	8072	240.0	253.000000	360.0	1.0	1	1	0	0	1	0	0
612	7583	0.0	187.000000	360.0	1.0	1	1	0	0	0	1	0
613	4583	0.0	133.000000	360.0	0.0	0	0	0	1	0	0	0

```
In [67]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score
from imblearn.over_sampling import SMOTE
```

```
In [69]: X=dataset.drop('Y',axis=1)
y=dataset['Y']
```

```
In [71]: print(X.shape)
print(X.isnull().sum())
```

```
(614, 14)
ApplicantIncome      0
CoapplicantIncome     0
LoanAmount            0
Loan_Amount_Term      0
Credit_History       0
Male                 0
Yes                  0
Not Graduate          0
Yes                  0
1                    0
2                    0
3+                   0
Semiurban             0
Urban                 0
dtype: int64
```

```
In [73]: print(y.shape)
print(y.isnull().sum())
```

```
(614,)
0
```

```
In [75]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)
```

```
In [77]: smote=SMOTE(random_state=42)
```

```
In [79]: X_train_balanced,y_train_balanced=smote.fit_resample(X_train,y_train)
```

```
In [81]: model=LogisticRegression(class_weight='balanced',max_iter=1000)
```

```
In [83]: model.fit(X_train_balanced,y_train_balanced)
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[83]: LogisticRegression
LogisticRegression(class_weight='balanced', max_iter=1000)
```

```
In [85]: y_pred=model.predict(X_test)
```

```
In [87]: accuracy=accuracy_score(y_test,y_pred)
```

```
In [89]: report=classification_report(y_test,y_pred)
```

```
In [91]: print("ACCURACY:",accuracy)
print("CLASSIFICATION:\n",report)
print("ROC-AUC Score",roc_auc_score(y_test,model.predict_proba(X_test)[:,-1]))
print("LOAN STATUS:",y_pred)
```

CLASSIFICATION:

ROC-AUC Score 0.750129132231405

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js